

CHENSHUHAO (CODY) QIN

Los Angeles, CA | (213) 922-0925 | cody2363264403@gmail.com | github.com/CodyQin | linkedin.com/in/chenshuhao-qin

EDUCATION

University of Southern California (USC) <i>Master of Science in Computer Science</i>	Los Angeles, CA <i>Jan. 2026 – Dec. 2027 (Expected)</i>
Coursework: Machine Learning, Deep Learning, Natural Language Processing, Data Structures & Algorithms, Operating Systems, Computer Architecture, Computer Networks, Predictive Modeling	
Duke University; Duke Kunshan University <i>Bachelor of Science in Computer Science; B.S. in Computation and Design</i>	Durham, NC / Kunshan, China <i>Aug. 2021 – May. 2025</i>
Honors: Apple WWDC23 Swift Student Challenge Winner (1 of 13 in China), Provincial Innovation Grant Recipient (\$3k), DKU Student Ambassador	

SKILLS

Programming: Python, TypeScript/JavaScript, Java, SQL, Swift, C
AI & Machine Learning: PyTorch, LLM Agents (AgentScope, LangChain, LlamaIndex), RAG, LoRA Fine-tuning, Ray, Hugging Face, TensorFlow, Pandas, NumPy
Infrastructure & Tools: Kubernetes, Docker, MySQL, PostgreSQL, Redis, FastAPI, React, Flask, CI/CD, Git, Linux, AWS

WORK EXPERIENCE

Alibaba Cloud <i>AI Engineer Intern</i>	Hangzhou, China <i>Jun. 2026 – Aug. 2026</i>
- Collaborated with Alibaba's Tongyi Lab to productize Data-Juicer Agents into natural-language-driven multimodal data-processing workflows for the broader Qwen ecosystem, replacing script-heavy pipelines on the Apsara Data Lake platform. - Built an agent-to-Ray distributed execution pipeline over private-cloud object storage and Kubernetes — 2,793 records, 16 operators, ~5 minutes — and landed the core capability upstream through a merged PR. - Designed a three-layer context-control architecture (tool-output compaction, memory compression, hard token budgets) cutting single-tool context payloads by 82% and enabling stable multi-turn agents on 30k-token models. - Shipped a React / FastAPI / MySQL agent workspace (persistent sessions, execution tracing, rollback, multi-instance isolation; 4 image releases) and contributed 5 merged PRs to the 7k+ Data-Juicer open-source ecosystem.	
Blueberry AI <i>AI Engineer Intern</i>	Burlingame, CA (Remote) <i>Nov. 2025 – Dec. 2025</i>
Built a Milvus + MySQL feature-extraction and auto-tagging pipeline batch-indexing 200K+ images for landmark / IP / SKU retrieval; benchmarked DINOv3, Qwen-VL, and BGE-VL under a unified recall / precision / latency framework for private deployment; delivered a YOLO + vector-search + Gradio retrieval POC validated for commercial use.	
Tsinghua Sichuan Energy Internet Research Institute <i>Algorithm Engineer Intern</i>	Chengdu, China <i>Sep. 2025 – Oct. 2025</i>
Engineered a DeepLabV3+ semantic-segmentation model for microgrid fault diagnosis reaching 92% accuracy across complex field scenarios; owned the full data lifecycle of 500+ annotated images from preprocessing to deployment, anchoring winning bids for major energy-infrastructure projects.	
Leiting Games (G-bits) <i>AI R&D Intern (AIGC & Automation)</i>	Shenzhen, China <i>May. 2025 – Aug. 2025</i>
Fine-tuned SDXL / Flux LoRA models on Kohya , shipping 30+ ad assets (80% adoption) that drove 117K+ installs and ¥24.6M billings (196% ROI) ; led a 5-person team scaling creative output 3× with 15–20% higher material ROI; built a 98.85%-accurate RPA monitor catching 1,000+ private-server violations (2025 Best Practice).	

PROJECTS

ChatDKU — RAG Campus Q&A Assistant <i>Python, LlamaIndex, LangChain, Redis</i>	Jul. 2024 – May. 2025
Co-founded a university-adopted RAG-agent Q&A system indexing 12,000+ webpages over a privately deployed LLM; implemented Hybrid Search (BM25 + vectors) , HyDE , and multi-level Redis caching, attaining ROUGE 81.1% / BLEU 60.2% ; authored the design as a DKU Signature Work.	

PUBLICATIONS

<i>Privacy-Preserving Room Occupancy Estimation Using Federated Analytics of BLE Packets — ACM SenSys 2024 (Demo Track, second author), pp. 889–890; 95% real-world accuracy via federated learning over BLE.</i>
